

Contents

Preface	ix
0 Initialization	1
0.1 An Effective Theory Approach	2
0.2 The Theoretical Minimum	3
1 Pretraining	11
1.1 Gaussian Integrals	12
1.2 Probability, Correlation and Statistics, and All That	21
1.3 Nearly-Gaussian Distributions	26
2 Neural Networks	37
2.1 Function Approximation	37
2.2 Activation Functions	43
2.3 Ensembles	47
3 Effective Theory of Deep Linear Networks at Initialization	53
3.1 Deep Linear Networks	54
3.2 Criticality	56
3.3 Fluctuations	59
3.4 Chaos	65
4 RG Flow of Preactivations	71
4.1 First Layer: Good-Old Gaussian	73
4.2 Second Layer: Genesis of Non-Gaussianity	79
4.3 Deeper Layers: Accumulation of Non-Gaussianity	90
4.4 Marginalization Rules	96
4.5 Subleading Corrections	100
4.6 RG Flow and RG Flow	103
5 Effective Theory of Preactivations at Initialization	109
5.1 Criticality Analysis of the Kernel	110
5.2 Criticality for Scale-Invariant Activations	123
5.3 Universality Beyond Scale-Invariant Activations	125

5.3.1 General Strategy 126

5.3.2 No Criticality: Sigmoid, Softplus, Nonlinear Monomials, etc. 128

5.3.3 $K^* = 0$ Universality Class: tanh, sin, etc. 130

5.3.4 Half-Stable Universality Classes: SWISH, etc. and GELU, etc. . . 135

5.4 Fluctuations 137

5.4.1 Fluctuations for the Scale-Invariant Universality Class 139

5.4.2 Fluctuations for the $K^* = 0$ Universality Class 141

5.5 Finite-Angle Analysis for the Scale-Invariant Universality Class 146

6 Bayesian Learning 153

6.1 Bayesian Probability 154

6.2 Bayesian Inference and Neural Networks 156

6.2.1 Bayesian Model Fitting 157

6.2.2 Bayesian Model Comparison 165

6.3 Bayesian Inference at Infinite Width 169

6.3.1 The Evidence for Criticality 169

6.3.2 Let’s Not Wire Together 173

6.3.3 Absence of Representation Learning 178

6.4 Bayesian Inference at Finite Width 179

6.4.1 Hebbian Learning, Inc. 179

6.4.2 Let’s Wire Together 182

6.4.3 Presence of Representation Learning 186

7 Gradient-Based Learning 191

7.1 Supervised Learning 192

7.2 Gradient Descent and Function Approximation 194

8 RG Flow of the Neural Tangent Kernel 199

8.0 Forward Equation for the NTK 200

8.1 First Layer: Deterministic NTK 206

8.2 Second Layer: Fluctuating NTK 207

8.3 Deeper Layers: Accumulation of NTK Fluctuations 211

8.3.0 *Interlude: Interlayer Correlations* 211

8.3.1 NTK Mean 215

8.3.2 NTK–Preactivation Cross Correlations 216

8.3.3 NTK Variance 221

9 Effective Theory of the NTK at Initialization 227

9.1 Criticality Analysis of the NTK 228

9.2 Scale-Invariant Universality Class 233

9.3 $K^* = 0$ Universality Class 236

9.4 Criticality, Exploding and Vanishing Problems, and None of That 241

10 Kernel Learning	247
10.1 A Small Step	248
10.1.1 No Wiring	250
10.1.2 No Representation Learning	250
10.2 A Giant Leap	252
10.2.1 Newton's Method	253
10.2.2 Algorithm Independence	257
10.2.3 <i>Aside</i> : Cross-Entropy Loss	259
10.2.4 Kernel Prediction	261
10.3 Generalization	264
10.3.1 Bias–Variance Tradeoff and Criticality	267
10.3.2 Interpolation and Extrapolation	277
10.4 Linear Models and Kernel Methods	282
10.4.1 Linear Models	282
10.4.2 Kernel Methods	284
10.4.3 Infinite-Width Networks as Linear Models	287
11 Representation Learning	291
11.1 Differential of the Neural Tangent Kernel	293
11.2 RG Flow of the dNTK	296
11.2.0 Forward Equation for the dNTK	297
11.2.1 First Layer: Zero dNTK	299
11.2.2 Second Layer: Nonzero dNTK	300
11.2.3 Deeper Layers: Growing dNTK	301
11.3 Effective Theory of the dNTK at Initialization	310
11.3.1 Scale-Invariant Universality Class	312
11.3.2 $K^* = 0$ Universality Class	314
11.4 Nonlinear Models and Nearly-Kernel Methods	317
11.4.1 Nonlinear Models	318
11.4.2 Nearly-Kernel Methods	324
11.4.3 Finite-Width Networks as Nonlinear Models	330
∞ The End of Training	335
$\infty.1$ Two More Differentials	337
$\infty.2$ Training at Finite Width	347
$\infty.2.1$ A Small Step Following a Giant Leap	351
$\infty.2.2$ Many Many Steps of Gradient Descent	358
$\infty.2.3$ Prediction at Finite Width	373
$\infty.3$ RG Flow of the ddNTKs: The Full Expressions	384
ϵ Epilogue: Model Complexity from the Macroscopic Perspective	389

A Information in Deep Learning **399**

A.1 Entropy and Mutual Information 400

A.2 Information at Infinite Width: Criticality 409

A.3 Information at Finite Width: Optimal Aspect Ratio 411

B Residual Learning **425**

B.1 Residual Multilayer Perceptrons 428

B.2 Residual Infinite Width: Criticality Analysis 429

B.3 Residual Finite Width: Optimal Aspect Ratio 431

B.4 Residual Building Blocks 436

References **439**

Index **445**