

# Denoising Diffusion Probabilistic Models and Generative Adversarial Networks

From VAEs and Diffusion Models to GANs — A Unified Derivation

Morten Hjorth-Jensen

Department of Physics & CCSE, University of Oslo

May 4-8, 2026

# Outline I

- 1 Motivation: VAEs and Their Limits
- 2 Hierarchical VAEs and the Diffusion Limit
- 3 The Forward Process
- 4 The Forward Posterior
- 5 The ELBO and Its Decomposition
- 6 Noise Prediction and the Training Objective
- 7 Connection to Score Matching
- 8 The Reverse Process and Sampling

# Outline II

- 9 Faster Sampling: DDIM
- 10 Latent Diffusion and Connections
- 11 Summary and Exercises
- 12 Generative Adversarial Networks
- 13 What Is a GAN?
- 14 Mathematical Formulation
- 15 Optimal Discriminator and JS Divergence
- 16 Loss Functions in Practice
- 17 Training Challenges

# Outline III

- 18 Evaluating GANs
- 19 Architecture Design
- 20 Extensions and Variants
- 21 Summary and Exercises
- 22 Comparing Generative Models: VAEs, Diffusion, and GANs

# Generative models: the common thread

All deep generative models share one goal: learn a distribution  $p_{\theta}(x)$  over data  $x$  that we can both evaluate and sample from.

| Family    | Latent                                         | Training signal          | Sample cost        |
|-----------|------------------------------------------------|--------------------------|--------------------|
| VAE       | $\mathbf{h} \in \mathbb{R}^{d_h}, d_h \ll d_x$ | ELBO                     | 1 decoder pass     |
| Diffusion | $\mathbf{x}_t \in \mathbb{R}^{d_x}$ (full dim) | <b>ELBO</b> , $T$ levels | $T$ denoiser calls |
| GAN       | $\mathbf{z} \in \mathbb{R}^k$                  | Adversarial              | 1 generator pass   |

## Key insight

Diffusion models are **hierarchical VAEs** with  $T$  latent levels, a fixed (non-learned) encoder, and full-dimensional latents. Understanding VAEs is the fastest route to understanding diffusion.

# The VAE: ELBO and reparameterisation (recap)

**Goal.** Maximise  $\log p_{\theta}(\mathbf{x})$  — intractable because  $p(\mathbf{x}) = \int p_{\theta}(\mathbf{x}|\mathbf{h}) p(\mathbf{h}) d\mathbf{h}$ .

Introduce approximate posterior  $q_{\phi}(\mathbf{h}|\mathbf{x})$ . By Jensen's inequality on  $\log$ :

$$\log p(\mathbf{x}) = \log \mathbb{E}_{q_{\phi}(\mathbf{h}|\mathbf{x})} \left[ \frac{p_{\theta}(\mathbf{x}, \mathbf{h})}{q_{\phi}(\mathbf{h}|\mathbf{x})} \right] \geq \mathbb{E}_{q_{\phi}(\mathbf{h}|\mathbf{x})} \left[ \log \frac{p_{\theta}(\mathbf{x}, \mathbf{h})}{q_{\phi}(\mathbf{h}|\mathbf{x})} \right] =: \mathcal{L}(\phi, \theta).$$

Splitting the joint  $p(\mathbf{x}, \mathbf{h}) = p_{\theta}(\mathbf{x}|\mathbf{h}) p(\mathbf{h})$ :

$$\boxed{\mathcal{L}(\phi, \theta) = \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{h}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{h})]}_{\text{reconstruction}} - \underbrace{D_{\text{KL}}(q_{\phi}(\mathbf{h}|\mathbf{x}) \parallel p(\mathbf{h}))}_{\text{regularisation}}}$$

**Reparameterisation.** With Gaussian encoder  $q_{\phi}(\mathbf{h}|\mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}_{\phi}, \boldsymbol{\sigma}_{\phi}^2 \mathbf{I})$  and prior  $p(\mathbf{h}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$ :

$$\mathbf{h} = \boldsymbol{\mu}_{\phi}(\mathbf{x}) + \boldsymbol{\sigma}_{\phi}(\mathbf{x}) \odot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}).$$

Closed-form KL:  $D_{\text{KL}}(q\|p) = \frac{1}{2} \sum_j \left( \mu_j^2 + \sigma_j^2 - \log \sigma_j^2 - 1 \right)$ .

# Why VAEs produce blurry samples

**The  $\ell_2$ -averaging problem.** The decoder minimises

$$-\mathbb{E}_{q_\phi}[\log p_\theta(\mathbf{x}|\mathbf{h})] \propto \mathbb{E}_{q_\phi}[\|\mathbf{x} - \hat{\mathbf{x}}_\theta(\mathbf{h})\|^2] .$$

Because the posterior  $q_\phi(\mathbf{h}|\mathbf{x})$  has nonzero variance,  $\hat{\mathbf{x}}_\theta$  learns the *mean over all likely completions*, which is systematically blurry.

## Other VAE limitations

- **Posterior collapse:** powerful decoder can ignore  $\mathbf{h}$ ,  $\text{KL} \rightarrow 0$ , latent useless.
- **ELBO gap:**  
 $\log p(\mathbf{x}) = \mathcal{L} + D_{\text{KL}}(q_\phi \| p(\mathbf{h}|\mathbf{x})) > \mathcal{L}.$
- Mean-field Gaussian approximation may miss multi-modal true posteriors.

## The diffusion solution

- Replace 1 latent level with  $T \gg 1$ .
- Fix the encoder: no collapse possible.
- Predict only the *noise added at one step*, not the full image — eliminates  $\ell_2$ -averaging artefacts.

# Markovian hierarchical VAE (MHVAE)

Generalise VAE to a **chain of  $T$  latent spaces**  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T$  (we already write  $\mathbf{x}_t$  for the  $t$ -th latent):

$$p(\mathbf{x}_0, \mathbf{x}_{1:T}) = p(\mathbf{x}_T) p_{\theta}(\mathbf{x}_0 | \mathbf{x}_1) \prod_{t=2}^T p_{\theta}(\mathbf{x}_{t-1} | \mathbf{x}_t) \quad (\text{generative})$$

$$q(\mathbf{x}_{1:T} | \mathbf{x}_0) = q(\mathbf{x}_1 | \mathbf{x}_0) \prod_{t=2}^T q(\mathbf{x}_t | \mathbf{x}_{t-1}) \quad (\text{inference})$$

Substituting into the ELBO and using the Markov property:

$$\log p(\mathbf{x}_0) \geq \mathbb{E}_{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \left[ \log \frac{p(\mathbf{x}_T) p_{\theta}(\mathbf{x}_0 | \mathbf{x}_1) \prod_{t=2}^T p_{\theta}(\mathbf{x}_{t-1} | \mathbf{x}_t)}{q(\mathbf{x}_1 | \mathbf{x}_0) \prod_{t=2}^T q(\mathbf{x}_t | \mathbf{x}_{t-1})} \right]$$

Each level contributes a reconstruction/regularisation pair.



# Three constraints that define a diffusion model

A **Variational Diffusion Model** (VDM; Kingma et al. 2021; Luo 2022) is an MHVAE with three additional constraints:

## Constraint 1 — Full-dimensional latents

$\dim(\mathbf{x}_t) = \dim(\mathbf{x}_0)$  for all  $t$ . No bottleneck; the latents live in *data space*.

## Constraint 2 — Fixed Gaussian encoder (forward process)

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{\alpha_t} \mathbf{x}_{t-1}, (1 - \alpha_t) \mathbf{I}), \quad \alpha_t = 1 - \beta_t \in (0, 1).$$

Not trained. The noise schedule  $\{\beta_t\}$  is a design choice.

## Constraint 3 — Gaussian terminal distribution

$\{\alpha_t\}$  is chosen so that  $q(\mathbf{x}_T | \mathbf{x}_0) \approx \mathcal{N}(\mathbf{0}, \mathbf{I})$  as  $t \rightarrow T$ . The model “forgets” the data completely.

## Forward process: the Gaussian product rule

**Key tool.** If  $\mathbf{x} = a\mathbf{u} + b\mathbf{v}$  with  $\mathbf{u} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  and  $\mathbf{v} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  independent, then

$$\mathbf{x} \sim \mathcal{N}(\mathbf{0}, (a^2 + b^2)\mathbf{I}).$$

**Proof of closed-form marginal.** Write one step as  $\mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}$  with  $\boldsymbol{\epsilon}_{t-1} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ . Unrolling two steps:

$$\begin{aligned}\mathbf{x}_t &= \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \boldsymbol{\epsilon}_{t-2}) + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1} \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \underbrace{\sqrt{\alpha_t (1 - \alpha_{t-1})} \boldsymbol{\epsilon}_{t-2} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}}_{\sim \mathcal{N}(\mathbf{0}, [\alpha_t(1 - \alpha_{t-1}) + (1 - \alpha_t)]\mathbf{I})} \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \bar{\boldsymbol{\epsilon}}\end{aligned}$$

since  $\alpha_t(1 - \alpha_{t-1}) + (1 - \alpha_t) = 1 - \alpha_t \alpha_{t-1}$ .

## Closed-form forward marginal $q(\mathbf{x}_t|\mathbf{x}_0)$

Define  $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$ . Continuing the induction:

### Closed-form marginal (key result)

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})$$

Equivalently, we can sample  $\mathbf{x}_t$  *directly* in one step:

$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}).$$

### Interpretation.

- $\sqrt{\bar{\alpha}_t}$  is the *signal coefficient*: decays toward 0.
- $\sqrt{1 - \bar{\alpha}_t}$  is the *noise coefficient*: grows toward 1.
- At  $t = 0$ :  $\mathbf{x}_0 = \mathbf{x}_0$  (clean).
- At  $t = T$ :  $\mathbf{x}_T \approx \boldsymbol{\epsilon}$  (pure noise).

No forward-pass through a network needed — this is just arithmetic.

# Signal-to-noise ratio and noise schedules

Define the **signal-to-noise ratio**:

$$\text{SNR}(t) = \frac{\bar{\alpha}_t}{1 - \bar{\alpha}_t}.$$

$\text{SNR}(0) \gg 1$  (clean data);  $\text{SNR}(T) \approx 0$  (pure noise). The VDM objective can be rewritten entirely in terms of  $\text{SNR}(t)$ , and the loss weight of each term equals  $\text{SNR}(t) - \text{SNR}(t + 1)$  (Kingma et al. 2021).

**Linear schedule** (Ho et al. 2020):

$$\beta_t = \beta_{\min} + \frac{t-1}{T-1}(\beta_{\max} - \beta_{\min})$$

$$\beta_{\min} = 10^{-4}, \beta_{\max} = 0.02, T = 1000.$$

Near-zero SNR early, then rapid collapse.

The cosine schedule avoids the very small  $\beta_t$  values that waste early training steps on near-identical images.

**Cosine schedule** (Nichol & Dhariwal 2021):

$$\bar{\alpha}_t = \frac{f(t)}{f(0)}, \quad f(t) = \cos^2\left(\frac{t/T + s}{1 + s} \cdot \frac{\pi}{2}\right)$$

$s = 0.008$ . SNR decreases more gradually, improving training at low  $t$ .

## Tractable forward posterior $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$

The *unconditional* reverse  $q(\mathbf{x}_{t-1}|\mathbf{x}_t)$  is intractable. But conditioned on  $\mathbf{x}_0$  it is tractable via Bayes' rule:

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0) q(\mathbf{x}_{t-1}|\mathbf{x}_0)}{q(\mathbf{x}_t|\mathbf{x}_0)} = \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1}) q(\mathbf{x}_{t-1}|\mathbf{x}_0)}{q(\mathbf{x}_t|\mathbf{x}_0)}.$$

The last equality uses the Markov property:  $q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0) = q(\mathbf{x}_t|\mathbf{x}_{t-1})$ .

Each factor is Gaussian:

$$\begin{aligned} q(\mathbf{x}_t|\mathbf{x}_{t-1}) &= \mathcal{N}(\sqrt{\alpha_t} \mathbf{x}_{t-1}, (1 - \alpha_t)\mathbf{I}) \\ q(\mathbf{x}_{t-1}|\mathbf{x}_0) &= \mathcal{N}(\sqrt{\bar{\alpha}_{t-1}} \mathbf{x}_0, (1 - \bar{\alpha}_{t-1})\mathbf{I}) \\ q(\mathbf{x}_t|\mathbf{x}_0) &= \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I}) \end{aligned}$$

Substituting and completing the square in  $\mathbf{x}_{t-1}$  yields a Gaussian in  $\mathbf{x}_{t-1}$ .

## Forward posterior: completing the square

The log of the numerator (as a function of  $\mathbf{x}_{t-1}$ ) is:

$$\begin{aligned} & -\frac{(\mathbf{x}_t - \sqrt{\alpha_t} \mathbf{x}_{t-1})^2}{2(1 - \alpha_t)} - \frac{(\mathbf{x}_{t-1} - \sqrt{\bar{\alpha}_{t-1}} \mathbf{x}_0)^2}{2(1 - \bar{\alpha}_{t-1})} \\ &= -\frac{1}{2} \left[ \frac{\alpha_t}{1 - \alpha_t} \mathbf{x}_{t-1}^2 - \frac{2\sqrt{\alpha_t} \mathbf{x}_t}{1 - \alpha_t} \mathbf{x}_{t-1} + \frac{1}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_{t-1}^2 - \frac{2\sqrt{\bar{\alpha}_{t-1}} \mathbf{x}_0}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_{t-1} + \text{const} \right] \end{aligned}$$

Grouping the  $\mathbf{x}_{t-1}^2$  and  $\mathbf{x}_{t-1}$  terms gives a precision:

$$\frac{1}{\tilde{\beta}_t} = \frac{\alpha_t}{1 - \alpha_t} + \frac{1}{1 - \bar{\alpha}_{t-1}} = \frac{\alpha_t(1 - \bar{\alpha}_{t-1}) + (1 - \alpha_t)}{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})} = \frac{1 - \bar{\alpha}_t}{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}.$$

Therefore:

$$\tilde{\beta}_t = \frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t.$$

## Forward posterior: the mean $\tilde{\boldsymbol{\mu}}_t$

The mean of  $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$  is read off from the linear term:

$$\frac{\tilde{\boldsymbol{\mu}}_t}{\tilde{\beta}_t} = \frac{\sqrt{\alpha_t} \mathbf{x}_t}{1 - \alpha_t} + \frac{\sqrt{\bar{\alpha}_{t-1}} \mathbf{x}_0}{1 - \bar{\alpha}_{t-1}}.$$

Multiplying through by  $\tilde{\beta}_t$ :

$$\tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0) = \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}} \beta_t}{1 - \bar{\alpha}_t} \mathbf{x}_0$$

## Forward posterior (exact result)

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I})$$

$$\tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0) = \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}} \beta_t}{1 - \bar{\alpha}_t} \mathbf{x}_0, \quad \tilde{\beta}_t = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$$

This is the *target* the learned reverse process must match.

Substituting the VDM structure into the MHVAE ELBO (Luo 2022, Sec. 3):

$$\begin{aligned} \log p(\mathbf{x}_0) \geq & \underbrace{\mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)}[\log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)]}_{\mathcal{L}_0: \text{reconstruction}} - \underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \| p(\mathbf{x}_T))}_{\mathcal{L}_T: \text{prior matching}} \\ & - \underbrace{\sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} \left[ D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \| p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)) \right]}_{\mathcal{L}_{1:T-1}: \text{denoising matching terms}} \end{aligned}$$

- $\mathcal{L}_T = 0$ : schedule ensures  $q(\mathbf{x}_T|\mathbf{x}_0) \approx \mathcal{N}(\mathbf{0}, \mathbf{I})$ .
- $\mathcal{L}_0$ : single reconstruction step, treated separately.
- $\mathcal{L}_{1:T-1}$ : dominant  $T - 1$  terms, each a Gaussian KL.



## Why condition on $\mathbf{x}_0$ reduces variance

**First attempt (naïve ELBO)** uses  $q(\mathbf{x}_{t-1}|\mathbf{x}_t)$  in the KL, which requires an expectation over *two* random variables per term. This has high Monte Carlo variance.

**Second attempt (conditioned ELBO)** uses  $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ . By the tower property:

$$\begin{aligned}\mathbb{E}_{q(\mathbf{x}_{t-1}, \mathbf{x}_t | \mathbf{x}_0)}[D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t))] \\ \geq \mathbb{E}_{q(\mathbf{x}_t | \mathbf{x}_0)}[D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t))].\end{aligned}$$

The conditioned form is a tighter bound *and* requires sampling only  $\mathbf{x}_t \sim q(\mathbf{x}_t|\mathbf{x}_0)$ , which is one-shot via the closed-form marginal. Each KL term now involves only a **single Monte Carlo draw**.

### Practical consequence

We compute  $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$  and evaluate the KL analytically — no expensive nested sampling.

## Gaussian KL: closed-form evaluation

Each denoising term is a KL between two Gaussians. For general Gaussians  $\mathcal{N}(\boldsymbol{\mu}_1, \Sigma_1)$  and  $\mathcal{N}(\boldsymbol{\mu}_2, \Sigma_2)$  in  $\mathbb{R}^d$ :

$$D_{\text{KL}}(\mathcal{N}(\boldsymbol{\mu}_1, \Sigma_1) \parallel \mathcal{N}(\boldsymbol{\mu}_2, \Sigma_2)) = \frac{1}{2} \left[ \log \frac{|\Sigma_2|}{|\Sigma_1|} - d + \text{tr}(\Sigma_2^{-1} \Sigma_1) + (\boldsymbol{\mu}_2 - \boldsymbol{\mu}_1)^\top \Sigma_2^{-1} (\boldsymbol{\mu}_2 - \boldsymbol{\mu}_1) \right].$$

In our case both distributions have variance  $\tilde{\beta}_t \mathbf{I}$ , so the  $\log |\cdot|$  and  $\text{tr}(\cdot)$  terms cancel exactly:

$$D_{\text{KL}}(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) \parallel p_{\boldsymbol{\theta}}(\mathbf{x}_{t-1} | \mathbf{x}_t)) = \frac{1}{2\tilde{\beta}_t} \|\tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0) - \boldsymbol{\mu}_{\boldsymbol{\theta}}(\mathbf{x}_t, t)\|^2 + C$$

where  $p_{\boldsymbol{\theta}}(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\boldsymbol{\mu}_{\boldsymbol{\theta}}(\mathbf{x}_t, t), \tilde{\beta}_t \mathbf{I})$  and  $C$  does not depend on  $\boldsymbol{\theta}$ .

**Training reduces to matching means.** The network only needs to predict  $\tilde{\boldsymbol{\mu}}_t$ .

## Reparameterising $\tilde{\mu}_t$ in terms of noise

From the closed-form forward marginal:  $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$ , we can solve for  $\mathbf{x}_0$ :

$$\mathbf{x}_0 = \frac{1}{\sqrt{\bar{\alpha}_t}} (\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \epsilon).$$

Substituting into  $\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0)$ :

$$\begin{aligned} \tilde{\mu}_t(\mathbf{x}_t, \epsilon) &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}} \beta_t}{1 - \bar{\alpha}_t} \cdot \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \epsilon}{\sqrt{\bar{\alpha}_t}} \\ &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1}) + \beta_t \sqrt{\bar{\alpha}_{t-1}/\bar{\alpha}_t}}{1 - \bar{\alpha}_t} \mathbf{x}_t - \frac{\beta_t \sqrt{\bar{\alpha}_{t-1}/\bar{\alpha}_t} \sqrt{1 - \bar{\alpha}_t}}{1 - \bar{\alpha}_t} \epsilon \end{aligned}$$

Using  $\bar{\alpha}_t = \bar{\alpha}_{t-1} \alpha_t$  so  $\sqrt{\bar{\alpha}_{t-1}/\bar{\alpha}_t} = 1/\sqrt{\alpha_t}$ , and collecting terms (see next slide):

# The noise-parameterised mean

After collecting terms in  $\mathbf{x}_t$  and  $\epsilon$ :

$$\tilde{\mu}_t = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon \right).$$

**Verification.** Numerator of  $\mathbf{x}_t$  coefficient:  $\alpha_t(1 - \bar{\alpha}_{t-1}) + \beta_t/\alpha_t \cdot \bar{\alpha}_{t-1}/\bar{\alpha}_{t-1} \dots$  simplifies to  $(1 - \bar{\alpha}_t)/\sqrt{\alpha_t}$ , giving coefficient  $1/\sqrt{\alpha_t}$ . ✓

If the network predicts noise  $\epsilon_{\theta}(\mathbf{x}_t, t) \approx \epsilon$ , then choose:

$$\mu_{\theta}(\mathbf{x}_t, t) = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right).$$

The denoising KL becomes:

$$\mathcal{L}_{t-1} = \mathbb{E}_{\mathbf{x}_0, \epsilon} \left[ \frac{\beta_t^2}{2\tilde{\beta}_t \alpha_t (1 - \bar{\alpha}_t)} \left\| \epsilon - \epsilon_{\theta}(\mathbf{x}_t, t) \right\|^2 \right]$$

where  $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$ .

# The simplified DDPM training objective

Ho et al. (2020) observed empirically that **dropping the time-dependent weighting**  $\beta_t^2 / (2\tilde{\beta}_t\alpha_t(1 - \bar{\alpha}_t))$  and averaging uniformly over  $t$  gives a simpler objective with *better* sample quality:

DDPM training loss  $\mathcal{L}_{\text{simple}}$

$$\mathcal{L}_{\text{simple}} = \mathbb{E}_{t \sim \mathcal{U}[1, T], \mathbf{x}_0, \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} \left[ \left\| \epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2 \right]$$

**Intuition.** Dropping the weight up-weights early steps (small  $t$ , small noise, easy to denoise) relative to later steps. This has a regularising effect: the network focuses more on the hard tasks where its predictions matter. The original weighted sum is still a valid ELBO; the simple loss is a reweighted ELBO that happens to work better in practice.

# Three equivalent parameterisations

$\mathbf{x}_t, \mathbf{x}_0, \epsilon$  satisfy a linear equation, so predicting any one is equivalent:

| Predict                                                                              | Objective                                                         | Notes                                 |
|--------------------------------------------------------------------------------------|-------------------------------------------------------------------|---------------------------------------|
| $\epsilon$ (noise)                                                                   | $\ \epsilon - \epsilon_{\theta}(\mathbf{x}_t, t)\ ^2$             | DDPM; most common                     |
| $\mathbf{x}_0$ (data)                                                                | $\ \mathbf{x}_0 - \hat{\mathbf{x}}_{\theta}(\mathbf{x}_t, t)\ ^2$ | Can over-smooth                       |
| $\mathbf{v} = \sqrt{\bar{\alpha}_t}\epsilon - \sqrt{1 - \bar{\alpha}_t}\mathbf{x}_0$ | $\ \mathbf{v} - \mathbf{v}_{\theta}(\mathbf{x}_t, t)\ ^2$         | Velocity; stable at $t \rightarrow 0$ |

**Conversions at inference:**

$$\hat{\mathbf{x}}_0 = \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_{\theta}}{\sqrt{\bar{\alpha}_t}}, \quad \hat{\epsilon} = \frac{\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \hat{\mathbf{x}}_0}{\sqrt{1 - \bar{\alpha}_t}}.$$

All three targets give the same denoising KL when substituted back.

# The score function

The **score** of a distribution  $p(\mathbf{x})$  is  $\nabla_{\mathbf{x}} \log p(\mathbf{x})$ . It points towards regions of high probability density.

**Score of the forward-diffused distribution.** Since  $q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})$ , its log-density is:

$$\log q(\mathbf{x}_t|\mathbf{x}_0) = -\frac{1}{2(1 - \bar{\alpha}_t)} \|\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \mathbf{x}_0\|^2 + C.$$

Differentiating with respect to  $\mathbf{x}_t$ :

$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}_0) = -\frac{\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \mathbf{x}_0}{1 - \bar{\alpha}_t} = -\frac{\boldsymbol{\epsilon}}{\sqrt{1 - \bar{\alpha}_t}}.$$

The score is proportional to the *negative noise*.

# DDPM as denoising score matching

Define the learned score network  $s_{\theta}(\mathbf{x}_t, t)$ . The **score matching objective** minimises:

$$\mathbb{E}_{t, \mathbf{x}_t} \left[ \left\| \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t) - s_{\theta}(\mathbf{x}_t, t) \right\|^2 \right].$$

The **denoising score matching** (Vincent 2011) substitution:

$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t) = \mathbb{E}_{q(\mathbf{x}_0|\mathbf{x}_t)} [\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t|\mathbf{x}_0)] = -\mathbb{E}_{q(\mathbf{x}_0|\mathbf{x}_t)} \left[ \frac{\epsilon}{\sqrt{1 - \bar{\alpha}_t}} \right].$$

With the optimal score network  $s_{\theta}^*(\mathbf{x}_t, t) = -\epsilon^*/\sqrt{1 - \bar{\alpha}_t}$ :

$$s_{\theta}(\mathbf{x}_t, t) = -\frac{\epsilon_{\theta}(\mathbf{x}_t, t)}{\sqrt{1 - \bar{\alpha}_t}} \implies s_{\theta}^* \approx \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t).$$

## Equivalence

**DDPM noise prediction = denoising score matching.** The DDPM network implicitly learns the score of the noisy data distribution at every noise level simultaneously.



## From score to SDE: the continuous-time limit

Song et al. (2021) embed the discrete chain into a continuous SDE. As  $T \rightarrow \infty$  with  $\beta_t \rightarrow \beta(t) dt$ , the forward process becomes:

$$d\mathbf{x} = -\frac{1}{2}\beta(t) \mathbf{x} dt + \sqrt{\beta(t)} d\mathbf{W}_t,$$

where  $\mathbf{W}_t$  is Brownian motion. By Anderson (1982), the **reverse SDE** is:

$$d\mathbf{x} = \left[ -\frac{1}{2}\beta(t) \mathbf{x} - \beta(t) \nabla_{\mathbf{x}} \log p_t(\mathbf{x}) \right] dt + \sqrt{\beta(t)} d\bar{\mathbf{W}}_t$$

where  $\bar{\mathbf{W}}_t$  is reverse-time Brownian motion. Replacing the score by the learned network gives:

$$d\mathbf{x} = \left[ -\frac{1}{2}\beta(t) \mathbf{x} + \beta(t) \frac{\epsilon_{\theta}(\mathbf{x}, t)}{\sqrt{1-\bar{\alpha}_t}} \right] dt + \sqrt{\beta(t)} d\bar{\mathbf{W}}_t.$$

The same network  $\epsilon_{\theta}$  that DDPM trains is the score estimator that drives the reverse SDE.

# The learned reverse process

Choose the learned reverse as:

$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_{\theta}(\mathbf{x}_t, t), \tilde{\beta}_t \mathbf{I})$$

with  $\boldsymbol{\mu}_{\theta}$  the noise-parameterised mean from Section 4.

**Why use  $\tilde{\beta}_t$  as the reverse variance?** The true posterior variance  $\tilde{\beta}_t$  lies between  $\beta_t$  (upper bound, for  $q(\mathbf{x}_{t-1}|\mathbf{x}_t)$ ) and  $\bar{\beta}_t$  (exact, for  $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ ). Ho et al. fixed it to  $\bar{\beta}_t$ ; later work (Nichol & Dhariwal 2021) learns it as an interpolation:

$$\Sigma_{\theta}(\mathbf{x}_t, t) = \exp(v \log \beta_t + (1 - v) \log \tilde{\beta}_t), \quad v = v_{\theta}(\mathbf{x}_t, t) \in [0, 1].$$

## DDPM reverse step

$$\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sqrt{\tilde{\beta}_t} \mathbf{z}, \quad \mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \ (t > 1), \ \mathbf{z} = \mathbf{0} \ (t = 1).$$

## Complete DDPM sampling procedure

- ❶ Sample  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ .
- ❷ For  $t = T, T - 1, \dots, 1$ :
  - ❶  $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ ;  $\mathbf{z} = \mathbf{0}$  if  $t = 1$ .
  - ❷ Predict noise:  $\hat{\epsilon} = \epsilon_{\theta}(\mathbf{x}_t, t)$ .
  - ❸ Update:  $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \hat{\epsilon} \right) + \sqrt{\tilde{\beta}_t} \mathbf{z}$ .
- ❸ Return  $\mathbf{x}_0$ .

- Each step is one network evaluation plus a cheap Gaussian update.
- With  $T = 1000$ : slow but high quality.
- **No discriminator**, no adversarial dynamics, no mode collapse.

# DDIM: deterministic ODE sampling

The DDPM score network is *independent of the sampling procedure*, so the stochastic SDE can be replaced by a deterministic ODE:

$$d\mathbf{x} = \left[ \frac{\mathbf{x} - \sqrt{1-\bar{\alpha}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}, \bar{\alpha})}{\sqrt{\bar{\alpha}}} \right] d\sqrt{\bar{\alpha}}.$$

Discretise over a subsequence  $\tau_1 < \dots < \tau_S$  ( $S \ll T$ ):

DDIM update ( $\sigma_{\tau_i} = \eta \sqrt{\tilde{\beta}_{\tau_i}}$ )

$$\begin{aligned} \mathbf{x}_{\tau_{i-1}} &= \sqrt{\bar{\alpha}_{\tau_{i-1}}} \hat{\mathbf{x}}_0 + \sqrt{1 - \bar{\alpha}_{\tau_{i-1}} - \sigma_{\tau_i}^2} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_{\tau_i}, \tau_i) + \sigma_{\tau_i} \mathbf{z}, \\ \hat{\mathbf{x}}_0 &= \frac{\mathbf{x}_{\tau_i} - \sqrt{1 - \bar{\alpha}_{\tau_i}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_{\tau_i}, \tau_i)}{\sqrt{\bar{\alpha}_{\tau_i}}}. \end{aligned}$$

- $\eta = 0$ : fully deterministic — same  $\mathbf{x}_T$  always gives same  $\mathbf{x}_0$ .
- $\eta = 1$ : recovers DDPM stochastic sampling.
- $S = 50$  matches quality of  $T = 1000$  DDPM steps.

# DDIM: derivation of the update rule

Start from the DDIM non-Markovian inference distribution (Song et al. 2020):

$$q_{\sigma}(\mathbf{x}_{\tau_{i-1}}|\mathbf{x}_{\tau_i}, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{\tau_{i-1}}; \sqrt{\bar{\alpha}_{\tau_{i-1}}} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_{\tau_{i-1}} - \sigma^2} \frac{\mathbf{x}_{\tau_i} - \sqrt{\bar{\alpha}_{\tau_i}} \mathbf{x}_0}{\sqrt{1 - \bar{\alpha}_{\tau_i}}}, \sigma^2 \mathbf{I}).$$

Replacing  $\mathbf{x}_0$  by the predicted  $\hat{\mathbf{x}}_0$  and the exact noise direction by  $\epsilon_{\theta}(\mathbf{x}_{\tau_i}, \tau_i)$  recovers the update rule on the previous slide.

**Why does this work?** The DDIM update is a first-order ODE solver for the probability-flow ODE. The ODE and the reverse SDE share the same marginals  $p_t(\mathbf{x}_t)$  at every time  $t$ . Since  $\epsilon_{\theta}$  approximates the score of  $p_t$ , integrating the ODE gives the same distribution over  $\mathbf{x}_0$  as ancestral sampling from the reverse SDE — but in far fewer steps.

# Latent Diffusion Models (Stable Diffusion)

**Problem.** DDPM in pixel space ( $512 \times 512 \times 3 = 786\,432$  dims) is too expensive.

**Solution** (Rombach et al. 2022): run diffusion in a compressed VAE latent space.

- ❶ **VAE encoder:**  $x \in \mathbb{R}^{H \times W \times 3} \rightarrow z \in \mathbb{R}^{H/8 \times W/8 \times 4}$  ( $\approx 48\times$  compression).
- ❷ **DDPM in  $z$ :** much cheaper than pixel space.
- ❸ **VAE decoder:**  $\hat{z}_0 \rightarrow \hat{x}$ .

## Why this works

- VAE latent is perceptually smooth and semantically structured.
- Diffusion on  $z$  skips low-level pixel correlations.
- Conditioning (text, depth, class) via cross-attention in the U-Net.
- KL regularisation keeps  $z$  smooth enough for the diffusion prior.

# VAE vs. Diffusion: mathematical parallel

Both maximise  $\log p(\mathbf{x}) \geq \mathcal{L}(\phi, \theta) = \mathbb{E}[\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{\text{KL}}(q_{\phi}||p)$ . Differences lie in how this framework is applied:

| Design axis      | VAE                            | DDPM                                |
|------------------|--------------------------------|-------------------------------------|
| Latent levels    | 1                              | $T = 300\text{--}1000$              |
| Encoder          | Learned $q_{\phi}$             | Fixed schedule                      |
| Bottleneck       | Yes ( $d_h \ll d_x$ )          | No ( $d_h = d_x$ )                  |
| Network predicts | $\mathbf{x}$ from $\mathbf{h}$ | $\epsilon$ from $(\mathbf{x}_t, t)$ |
| KL terms         | 1                              | $T - 1$                             |
| Training loss    | BCE + $D_{\text{KL}}$          | MSE (noise)                         |
| Sampling cost    | 1 decoder call                 | $T$ denoiser calls                  |

## Key insight

DDPM = VAE with (a)  $T$  latent levels, (b) fixed encoder, (c) no bottleneck.

# VAE vs. Diffusion: pros and cons

## VAE

### Pros:

- Fast inference (1 pass)
- Compact structured latent space
- Stable training, closed-form KL
- Good for clustering / interpolation
- Scalable to large datasets

### Cons:

- **Blurry samples** ( $\ell_2$ -averaging)
- Posterior collapse risk
- ELBO gap (log-likelihood not exact)
- Mean-field limits posterior quality

## Diffusion (DDPM)

### Pros:

- **State-of-the-art** sample quality
- No mode collapse
- Stable training (MSE, no adversary)
- Grounded via ELBO + score matching
- Flexible conditioning

### Cons:

- **Slow** sampling ( $T$  evaluations)
- No compact latent space
- High memory for large  $T$
- Less natural for representation learning



## Key equations: complete reference

| Concept                | Formula                                                                                                                                                                                   |
|------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Forward step           | $q(\mathbf{x}_t \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{\alpha_t} \mathbf{x}_{t-1}, (1 - \alpha_t)\mathbf{I})$                                                                              |
| Forward marginal       | $q(\mathbf{x}_t \mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})$                                                                          |
| One-shot sample        | $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}$                                                                                     |
| Forward posterior mean | $\tilde{\boldsymbol{\mu}}_t = \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} \mathbf{x}_0$ |
| Forward posterior var. | $\tilde{\beta}_t = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$                                                                                                             |
| Noise-param. mean      | $\tilde{\boldsymbol{\mu}}_t = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon} \right)$                                    |
| Training loss          | $\mathcal{L}_{\text{simple}} = \mathbb{E}[\ \boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t)\ ^2]$                                                                 |
| Score relation         | $\mathbf{s}_{\theta}(\mathbf{x}_t, t) = -\boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) / \sqrt{1 - \bar{\alpha}_t}$                                                                     |
| Reverse step           | $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta} \right) + \sqrt{\tilde{\beta}_t} \mathbf{z}$ |
| SNR                    | $\text{SNR}(t) = \bar{\alpha}_t / (1 - \bar{\alpha}_t)$                                                                                                                                   |

## Exercises 1–3

- 1 **Forward marginal.** Prove  $q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})$  by induction from  $q(\mathbf{x}_t|\mathbf{x}_{t-1})$ .
- 2 **Forward posterior.** Apply Bayes' rule to derive  $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ . Show that completing the square gives  $\tilde{\boldsymbol{\mu}}_t$  and  $\tilde{\beta}_t$ .
- 3 **Gaussian KL.** Derive  $D_{\text{KL}}(\mathcal{N}(\boldsymbol{\mu}_1, \sigma^2\mathbf{I})\|\mathcal{N}(\boldsymbol{\mu}_2, \sigma^2\mathbf{I})) = \|\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2\|^2/(2\sigma^2)$  and verify that the trace and log-det terms cancel exactly.

## Exercises 4–6

- 4 **Noise reparameterisation.** Substitute  $\mathbf{x}_0 = (\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t}\boldsymbol{\epsilon})/\sqrt{\bar{\alpha}_t}$  into  $\tilde{\boldsymbol{\mu}}_t$  and verify the result equals  $\frac{1}{\sqrt{\alpha_t}}(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}}\boldsymbol{\epsilon})$ .
- 5 **Score matching.** Write the score of  $q(\mathbf{x}_t|\mathbf{x}_0)$  in terms of  $\boldsymbol{\epsilon}$  and show the DDPM loss equals denoising score matching.
- 6 **ELBO decomposition.** Derive the three-term decomposition  $\mathcal{L}_0 + \mathcal{L}_T + \sum_t \mathcal{L}_{t-1}$  and explain why  $\mathcal{L}_T = 0$  by design.

- ① D. Kingma & M. Welling, *Auto-Encoding Variational Bayes*, ICLR 2014. arXiv:1312.6114
- ② J. Sohl-Dickstein et al., *Deep Unsupervised Learning using Nonequilibrium Thermodynamics*, ICML 2015. arXiv:1503.03585
- ③ J. Ho, A. Jain & P. Abbeel, *Denoising Diffusion Probabilistic Models*, NeurIPS 2020. arXiv:2006.11239
- ④ J. Song, C. Meng & S. Ermon, *Denoising Diffusion Implicit Models*, ICLR 2021. arXiv:2010.02502
- ⑤ Y. Song et al., *Score-Based Generative Modeling through Stochastic Differential Equations*, ICLR 2021. arXiv:2011.13456
- ⑥ D. Kingma et al., *Variational Diffusion Models*, NeurIPS 2021. arXiv:2107.00630
- ⑦ A. Nichol & P. Dhariwal, *Improved Denoising Diffusion Probabilistic Models*, ICML 2021. arXiv:2102.09672
- ⑧ R. Rombach et al., *High-Resolution Image Synthesis with Latent Diffusion Models*, CVPR 2022. arXiv:2112.10752
- ⑨ C. Luo, *Understanding Diffusion Models: A Unified Perspective*, 2022. arXiv:2208.11970

## Transition: From Likelihood to Adversarial Training

So far we studied models maximising (a lower bound on)  $\log p(\mathbf{x})$ :

| Model      | Objective                   | Key idea                              |
|------------|-----------------------------|---------------------------------------|
| VAE        | Maximise ELBO               | Variational inference, 1 latent level |
| DDPM       | Maximise ELBO ( $T$ levels) | Fixed encoder, noise prediction       |
| <b>GAN</b> | <b>No likelihood bound</b>  | <b>Adversarial game</b>               |

### Fundamental difference

GANs never evaluate  $\log p_{\theta}(\mathbf{x})$ . A **discriminator**  $D$  provides the training signal by classifying real vs. fake, driving the **generator**  $G$  to match  $p_r$  — without any likelihood computation.

*Benefit:* sharper samples, no  $\ell_2$ -blurring.    *Cost:* no likelihood, mode collapse risk, instability.

# Generative Adversarial Networks

A **GAN** (Goodfellow et al., NeurIPS 2014) pits two neural networks against each other in a zero-sum game:

## Generator $G$

- Input: noise  $z \sim p_z(z)$
- Output: synthetic sample  $x = G(z; \theta^{(g)})$
- Goal: fool the discriminator
- Implicitly defines distribution  $p_g$


## Discriminator $D$

- Input: image  $x$
- Output:  $D(x; \theta^{(d)}) \in (0, 1)$   
 $= \Pr(x \text{ is real})$
- Goal: distinguish real from fake

## Key insight

Neither network sees the other's parameters. They improve *each other* through adversarial feedback alone.

# GAN Architecture (Schematic)



- $G$  maps  $z \in \mathbb{R}^k \rightarrow \mathbb{R}^d$  (noise  $\rightarrow$  data space)
- $D$  maps  $x \in \mathbb{R}^d \rightarrow (0, 1)$  (image  $\rightarrow$  probability)
- Training alternates: freeze  $G$  and update  $D$ ; freeze  $D$  and update  $G$

## Three learning settings:

- 1 **Unsupervised** — no labels;  $G$  learns the full data manifold
- 2 **Supervised** — conditional GAN conditions  $G$  on class label
- 3 **Semi-supervised** —  $D$  doubles as a feature extractor

## Applications:

- Image synthesis / super-resolution
- Text-to-image translation
- Video prediction
- Data augmentation for rare events
- Anomaly detection
- Scientific simulation:  
particle physics, drug discovery,  
climate modelling

# Notation and Probability Spaces

| Symbol                                                | Meaning                                                               |
|-------------------------------------------------------|-----------------------------------------------------------------------|
| $p_r(\mathbf{x})$                                     | Real data distribution (unknown; represented by training set)         |
| $p_z(\mathbf{z})$                                     | Latent noise prior, e.g. $\mathcal{N}(\mathbf{0}, I)$                 |
| $G(\mathbf{z}; \boldsymbol{\theta}^{(g)})$            | Generator: differentiable map $\mathbb{R}^k \rightarrow \mathbb{R}^d$ |
| $p_g$                                                 | Generator distribution = push-forward of $p_z$ through $G$            |
| $D(\mathbf{x}; \boldsymbol{\theta}^{(d)}) \in (0, 1)$ | Discriminator: $\Pr(\mathbf{x} \text{ is real})$                      |

**Implicit distribution.** For invertible  $G$  the change-of-variables gives:

$$p_g(\mathbf{x}) = p_z(G^{-1}(\mathbf{x})) \left| \det \frac{\partial G^{-1}}{\partial \mathbf{x}} \right|.$$

In practice  $G$  is not invertible;  $p_g$  is an *implicit* distribution accessed only through sampling.



# The Minimax Objective

## Value function (Goodfellow et al., 2014)

$$\min_G \max_D V(G, D) = \underbrace{\mathbb{E}_{\mathbf{x} \sim p_r} [\log D(\mathbf{x})]}_{D \text{ correct on real}} + \underbrace{\mathbb{E}_{\mathbf{z} \sim p_z} [\log(1 - D(G(\mathbf{z})))]}_{D \text{ correct on fake}}$$

### Discriminator (maximise $V$ ):

- $D(\mathbf{x}) \rightarrow 1$  on real data
- $D(G(\mathbf{z})) \rightarrow 0$  on generated data

### Generator (minimise $V$ ):

- $D(G(\mathbf{z})) \rightarrow 1$  (fool  $D$ )
- $\mathbb{E}_{p_r} [\log D(\mathbf{x})]$  is constant w.r.t.  $G$ , so  $G$  minimises  $\mathbb{E}[\log(1 - D(G(\mathbf{z})))]$

# Non-Saturating Generator Loss

**Problem with the original objective.** Early in training  $D$  is strong, so  $D(G(\mathbf{z})) \approx 0$ :

$$\frac{\partial}{\partial \theta^{(g)}} \log(1 - D(G(\mathbf{z}))) \approx \frac{\partial}{\partial \theta^{(g)}} \log 1 = 0.$$

The gradient *vanishes* before  $G$  has learned anything.

## Fix: non-saturating loss

Instead of minimising  $\log(1 - D(G(\mathbf{z})))$ , maximise  $\log D(G(\mathbf{z}))$ :

$$\mathcal{L}_G^{\text{non-sat}} = -\mathbb{E}_{\mathbf{z} \sim p_z} [\log D(G(\mathbf{z}))].$$

Same Nash equilibrium; stronger gradient when  $D(G(\mathbf{z})) \approx 0$ .

| Loss form               | Value at $D(G(\mathbf{z})) = \varepsilon \approx 0$ | gradient                                   |
|-------------------------|-----------------------------------------------------|--------------------------------------------|
| $\log(1 - \varepsilon)$ | $\approx -\varepsilon$                              | $\approx 1$ (small)                        |
| $-\log \varepsilon$     | $+\infty$                                           | $1/\varepsilon \rightarrow \infty$ (large) |

# The Alternating Gradient Algorithm

## One training iteration

- 1 Freeze  $G$ ; gradient ascent on  $D$ :

$$\theta^{(d)} \leftarrow \theta^{(d)} + \eta \nabla_{\theta^{(d)}} V(G, D).$$

- 2 Freeze  $D$ ; gradient descent on  $G$ :

$$\theta^{(g)} \leftarrow \theta^{(g)} - \eta \nabla_{\theta^{(g)}} \mathcal{L}_G.$$

- One  $D$  step per  $G$  step in practice ( $k = 1$ ).
- Fake images are **detached** from the computation graph during the  $D$  step to prevent spurious gradients flowing into  $G$ .
- Both networks update simultaneously — this does *not* guarantee convergence (see Section ??).

# Deriving the Optimal Discriminator

Fix  $G$  (hence fix  $p_g$ ). Write  $V$  as an integral:

$$V(G, D) = \int_{\mathbf{x}} [p_r(\mathbf{x}) \log D(\mathbf{x}) + p_g(\mathbf{x}) \log(1 - D(\mathbf{x}))] d\mathbf{x}.$$

For each  $\mathbf{x}$  independently, maximise the integrand  $f(t) = A \log t + B \log(1 - t)$ ,  $A = p_r(\mathbf{x})$ ,  $B = p_g(\mathbf{x})$ ,  $t = D(\mathbf{x}) \in (0, 1)$ :

$$f'(t) = \frac{A}{t} - \frac{B}{1-t} = \frac{A - (A+B)t}{t(1-t)} = 0 \implies t^* = \frac{A}{A+B}.$$

Check:  $f''(t^*) = -A/(t^*)^2 - B/(1-t^*)^2 < 0$  ✓ (maximum).

## Optimal discriminator

$$D^*(\mathbf{x}) = \frac{p_r(\mathbf{x})}{p_r(\mathbf{x}) + p_g(\mathbf{x})}.$$

This is the **Bayes-optimal classifier** for the two-class problem {real, fake} under equal class priors.

# Nash Equilibrium and Loss at the Optimum

**Nash equilibrium.** At  $p_g = p_r$ :

$$D^*(\mathbf{x}) = \frac{p_r(\mathbf{x})}{p_r(\mathbf{x}) + p_r(\mathbf{x})} = \frac{1}{2}.$$

The discriminator cannot distinguish real from fake — it is reduced to random guessing.

**Loss at the equilibrium.** Substituting  $D^* = 1/2$ :

$$\begin{aligned} V(G^*, D^*) &= \log \frac{1}{2} \int p_r d\mathbf{x} + \log \frac{1}{2} \int p_r d\mathbf{x} \\ &= -2 \log 2 \approx -1.386. \end{aligned}$$

## Convergence indicator

In practice, monitor  $D(\mathbf{x})$  and  $D(G(\mathbf{z}))$  during training. Both should converge to 0.5 at equilibrium. If  $D(G(\mathbf{z})) \ll 0.5$ , the generator is losing. If  $D(\mathbf{x}) \ll 0.5$ , the discriminator is failing.

# KL and JS Divergences

**Kullback-Leibler divergence:**

$$D_{\text{KL}}(p\|q) = \int p(\mathbf{x}) \log \frac{p(\mathbf{x})}{q(\mathbf{x})} d\mathbf{x} \geq 0,$$

with equality iff  $p = q$  a.e. (Gibbs' inequality). Note:  $D_{\text{KL}}(p\|q) \neq D_{\text{KL}}(q\|p)$  in general.

**Jensen-Shannon divergence** (symmetric, bounded):

$$D_{\text{JS}}(p\|q) = \frac{1}{2} D_{\text{KL}}\left(p \parallel \frac{p+q}{2}\right) + \frac{1}{2} D_{\text{KL}}\left(q \parallel \frac{p+q}{2}\right).$$

## Properties of $D_{\text{JS}}$

- Symmetric:  $D_{\text{JS}}(p\|q) = D_{\text{JS}}(q\|p)$
- Non-negative:  $D_{\text{JS}}(p\|q) \geq 0$ , with equality iff  $p = q$
- Bounded:  $D_{\text{JS}}(p\|q) \in [0, \log 2]$  for any  $p, q$
- Related to total variation:  $D_{\text{JS}}(p\|q) \leq \log 2 \cdot \text{TV}(p, q)$

# GAN Loss = Jensen-Shannon Divergence

Expanding  $D_{\text{JS}}(p_r \| p_g)$  and simplifying:

$$\begin{aligned} D_{\text{JS}}(p_r \| p_g) &= \frac{1}{2} D_{\text{KL}}\left(p_r \left\| \frac{p_r + p_g}{2}\right.\right) + \frac{1}{2} D_{\text{KL}}\left(p_g \left\| \frac{p_r + p_g}{2}\right.\right) \\ &= \frac{1}{2} \left( \log 2 + \int p_r \log \frac{p_r}{p_r + p_g} d\mathbf{x} \right) + \frac{1}{2} \left( \log 2 + \int p_g \log \frac{p_g}{p_r + p_g} d\mathbf{x} \right). \end{aligned}$$

Recognising  $V(G, D^*)$  in the integrals:

Main theorem (Goodfellow et al. 2014)

$$V(G, D^*) = 2 D_{\text{JS}}(p_r \| p_g) - 2 \log 2.$$

**Conclusion.** The GAN minimax objective is equivalent to minimising the Jensen-Shannon divergence between  $p_r$  and  $p_g$ . The global minimum  $-2 \log 2$  is achieved when  $p_g = p_r$ .

# Binary Cross-Entropy Losses

**Discriminator loss** (real label =  $s$ , fake label = 0):

$$\mathcal{L}_D = -\frac{1}{N} \sum_{i=1}^N \log D(\mathbf{x}_i) - \frac{1}{M} \sum_{j=1}^M \log(1 - D(G(\mathbf{z}_j))).$$

**Generator loss** (non-saturating):

$$\mathcal{L}_G = -\frac{1}{M} \sum_{j=1}^M \log D(G(\mathbf{z}_j)).$$

## One-sided label smoothing ( $s = 0.9$ )

Replace real target 1 with  $s < 1$ . The optimal discriminator becomes

$D_s^*(\mathbf{x}) = s \cdot p_r(\mathbf{x}) / (p_r(\mathbf{x}) + p_g(\mathbf{x}))$ , preventing overconfidence and keeping gradients to  $G$  alive.



# Wasserstein Distance (WGAN)

**Problem with JS divergence.** When  $p_r$  and  $p_g$  have non-overlapping supports (common early in training):  $D_{JS}(p_r \| p_g) = \log 2$  constantly, giving zero gradient everywhere.

**Wasserstein-1 distance** (*Earth-mover*):

$$W(p_r, p_g) = \inf_{\gamma \in \Pi(p_r, p_g)} \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim \gamma} \|\mathbf{x} - \mathbf{y}\|$$

$\Pi(p_r, p_g)$ : all couplings with marginals  $p_r, p_g$ .

## Kantorovich-Rubinstein duality

$$W(p_r, p_g) = \sup_{\|f\|_L \leq 1} (\mathbb{E}_{p_r}[f] - \mathbb{E}_{p_g}[f]).$$

Sup. over all 1-Lipschitz functions  $f$ .

## Advantage

$W$  gives meaningful gradients even when  $p_r \cap p_g = \emptyset$ . Training is more stable.

# WGAN: Gradient Penalty

**WGAN critic.** Replace  $D$  with a critic  $f_w$  (no sigmoid):

$$\mathcal{L}_{\text{critic}} = \mathbb{E}_{p_g}[f_w(\mathbf{x})] - \mathbb{E}_{p_r}[f_w(\mathbf{x})] + \lambda \mathcal{L}_{\text{GP}}.$$

**Gradient penalty** (Gulrajani et al., 2017):

$$\mathcal{L}_{\text{GP}} = \mathbb{E}_{\hat{\mathbf{x}}} \left[ \left( \|\nabla_{\hat{\mathbf{x}}} f_w(\hat{\mathbf{x}})\|_2 - 1 \right)^2 \right],$$

where  $\hat{\mathbf{x}} = \epsilon \mathbf{x} + (1 - \epsilon)G(\mathbf{z})$ ,  $\epsilon \sim U[0, 1]$  (uniform on straight lines between real and fake points).

- Enforces the 1-Lipschitz constraint via soft penalisation.
- More stable than weight clipping (original WGAN).
- $\lambda = 10$  is a standard choice.

**When does it occur?** When  $D$  is perfect early in training:  $D(\mathbf{x}) = 1 \ \forall \mathbf{x} \sim p_r$  and  $D(G(\mathbf{z})) = 0 \ \forall \mathbf{z} \sim p_z$ .

**Original  $G$  gradient:**

$$\nabla_{\theta(g)} \mathbb{E}[\log(1 - D(G(\mathbf{z}))) ] \approx \nabla_{\theta(g)} \mathbb{E}[\log 1] = \mathbf{0}.$$

No information flows to  $G$ .

**Non-saturating loss  $-\log D(G(\mathbf{z}))$ :**

$$\nabla_{\theta(g)} (-\log D(G(\mathbf{z}))) = -\frac{1}{D(G(\mathbf{z}))} \nabla_{\theta(g)} D(G(\mathbf{z})).$$

When  $D(G(\mathbf{z})) \approx 0$ , the factor  $1/D(G(\mathbf{z})) \rightarrow \infty$  amplifies the gradient — learning continues.

**Definition.** The generator maps many different noise vectors  $z$  to the *same* (or very few) output(s):

$$\text{supp}(p_g) \ll \text{supp}(p_r).$$

The JS divergence can still be low on the covered modes, so the loss does not detect collapse.

**Why it happens.** Given the current discriminator,  $G$  finds a single “safe” output that fools  $D$ ;  $D$  then adapts, and  $G$  shifts to another single output. The two networks chase each other indefinitely.

**Mitigations:**

- **Minibatch discrimination:**  $D$  sees statistics of a batch, not individual images — diversity is rewarded.
- **Unrolled GANs:**  $G$  optimises against a future  $D$ , preventing short-sighted collapse.
- **Spectral normalisation:** stabilises  $D$  by bounding its Lipschitz constant.
- **WGAN:** Wasserstein distance measures full distribution mismatch, not just modes.

# Non-Convergence and the Game-Theoretic View

GAN training solves a **two-player non-cooperative game**:

$$G^* = \arg \min_G \arg \max_D V(G, D).$$

**Nash equilibrium** exists at  $p_g = p_r$ ,  $D^* \equiv \frac{1}{2}$ , but:

- Simultaneous gradient updates do *not* guarantee convergence to a Nash equilibrium.
- The loss landscape of each player changes at every step (non-stationary environment).
- If  $\max_D V(G, D)$  is not convex in  $\theta^{(g)}$ , global convergence fails even in theory.

**Practical stabilisation:**

- Lower  $\beta_1$  in Adam (reduces momentum: 0.5 instead of 0.9)
- One-sided label smoothing
- Spectral normalisation of  $D$ 's weights
- Progressive growing (ProGAN), self-attention (SAGAN)
- Two-time-scale updates: different learning rates for  $G$  and  $D$

# Evaluation Metrics: Overview

Evaluating generative models is hard: there is no ground truth output to compare against.

| Metric           |       | Measures                                         | Limitation                     |
|------------------|-------|--------------------------------------------------|--------------------------------|
| Inception (IS)   | Score | Quality + diversity of generated images          | Requires pretrained classifier |
| FID              |       | Distance between real/fake feature distributions | Sensitive to sample size       |
| MMD              |       | Kernel-based distribution distance               | Choice of kernel matters       |
| Precision/Recall |       | Coverage vs. fidelity separately                 | Requires feature extractor     |

# Maximum Mean Discrepancy (MMD)

**Definition in RKHS.** Let  $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  be a kernel. The **MMD** between distributions  $p, q$  is:

$$\text{MMD}^2(p, q) = \mathbb{E}_{\mathbf{x}, \mathbf{x}' \sim p}[k(\mathbf{x}, \mathbf{x}')] - 2 \mathbb{E}_{\mathbf{x} \sim p, \mathbf{y} \sim q}[k(\mathbf{x}, \mathbf{y})] + \mathbb{E}_{\mathbf{y}, \mathbf{y}' \sim q}[k(\mathbf{y}, \mathbf{y}')].$$

With the Gaussian kernel  $k(\mathbf{x}, \mathbf{y}) = \exp(-\|\mathbf{x} - \mathbf{y}\|^2 / 2\sigma^2)$ : MMD = 0 iff  $p = q$  (characteristic kernel property).

## Unbiased empirical estimator

Given  $n$  real samples  $\{\mathbf{x}_i\}$  and  $m$  fake samples  $\{\mathbf{y}_j\}$ :

$$\widehat{\text{MMD}}^2 = \frac{1}{n(n-1)} \sum_{i \neq j} k(\mathbf{x}_i, \mathbf{x}_j) - \frac{2}{nm} \sum_{i,j} k(\mathbf{x}_i, \mathbf{y}_j) + \frac{1}{m(m-1)} \sum_{i \neq j} k(\mathbf{y}_i, \mathbf{y}_j).$$

Diagonal terms ( $i = j$ ) excluded for unbiasedness.

# Fréchet Inception Distance (FID)

**FID** is currently the most widely used GAN metric.

## Procedure:

- 1 Extract feature vectors from the **InceptionV3** penultimate layer for  $n$  real images and  $n$  generated images.
- 2 Fit Gaussians:  $\mathcal{N}(\mu_r, \Sigma_r)$  and  $\mathcal{N}(\mu_g, \Sigma_g)$ .
- 3 Compute the **Fréchet distance**:

$$\text{FID} = \|\mu_r - \mu_g\|^2 + \text{Tr}\left(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{1/2}\right).$$

- Lower FID = better quality and diversity.
- State-of-the-art GANs achieve  $\text{FID} < 5$  on MNIST;  $< 10$  on CIFAR-10.
- Sensitive to sample size  $n$  — use  $\geq 10\,000$  samples.



# Generator Architecture for MNIST

**Fully-connected generator** (Goodfellow et al., 2014 style):

$$z \in \mathbb{R}^{100} \xrightarrow{\text{Lin}+\text{BN}+\text{LReLU}} \mathbb{R}^{256} \xrightarrow{\text{Lin}+\text{BN}+\text{LReLU}} \mathbb{R}^{512} \xrightarrow{\text{Lin}+\text{BN}+\text{LReLU}} \mathbb{R}^{1024} \xrightarrow{\text{Lin}+\text{Tanh}} \mathbb{R}^{784}$$

**Design choices:**

- **BatchNorm** after each Linear layer: normalises activations, allows higher LR, stabilises gradient flow
- **LeakyReLU(0.2)**: prevents dying ReLU; small gradient for negative inputs
- **Tanh** output: maps to  $[-1, 1]$ , matching MNIST normalised to  $[-1, 1]$
- **Weights**:  $\mathcal{N}(0, 0.02^2)$  (DCGAN convention)

**Discriminator differences:**

- **No BatchNorm**: introduces minibatch dependencies, causes instability in  $D$
- **Dropout(0.3)**: regularises, prevents  $D$  from over-fitting on real data
- **Sigmoid** output: probability in  $(0, 1)$
- **Input**: flattened image  $\in \mathbb{R}^{784}$

# Latent Space Interpolation

A well-trained generator has a **smooth latent space**. Linear interpolation between two noise vectors:

$$z(\alpha) = (1 - \alpha) z_1 + \alpha z_2, \quad \alpha \in [0, 1]$$

should yield a smooth visual transition between the corresponding images.

- **Smooth transitions:** generator has learned a good manifold.
- **Abrupt jumps:** mode collapse or poor latent coverage.
- Provides a qualitative test of latent space geometry that FID cannot capture.

## Spherical interpolation

On a high-dimensional sphere (approximately the case for  $z \sim \mathcal{N}(0, I)$ ), *spherical* interpolation  $z(\alpha) = \frac{\sin((1-\alpha)\Omega)}{\sin \Omega} z_1 + \frac{\sin(\alpha\Omega)}{\sin \Omega} z_2$ ,  $\cos \Omega = \hat{z}_1 \cdot \hat{z}_2$ , stays on the noise manifold and gives more natural transitions.

# Conditional GAN (cGAN)

A **conditional GAN** (Mirza & Osindero, 2014) conditions both  $G$  and  $D$  on auxiliary information  $\mathbf{y}$  (class label, image, text):

$$\min_G \max_D V(G, D) = \mathbb{E}[\log D(\mathbf{x}|\mathbf{y})] \\ + \mathbb{E}[\log(1 - D(G(\mathbf{z}|\mathbf{y})))].$$

For MNIST:  $\mathbf{y}$  is the digit class (one-hot,  $\mathbb{R}^{10}$ ).  
Concatenate  $\mathbf{y}$  to both  $G$ 's input  $\mathbf{z}$  and  $D$ 's input  $\mathbf{x}$ .

## cGAN enables

- Controlled generation: specify which digit to generate
- Image-to-image translation (pix2pix)
- Text-to-image synthesis
- Style transfer

# Deep Convolutional GAN (DCGAN)

Radford, Metz, Chintala (ICLR 2016) replaced fully-connected layers with **strided convolutions**:

- **Generator**: fractionally strided convolutions (transposed conv) to upsample from  $z \in \mathbb{R}^{100}$  to full-resolution image.
- **Discriminator**: strided convolutions to downsample.
- BatchNorm in  $G$  (all layers except output); BatchNorm in  $D$  (except input).
- ReLU in  $G$ ; LeakyReLU in  $D$ .
- No pooling layers; no fully-connected layers (except for  $z$ ).

## Key advance

DCGAN produced the first visually convincing high-resolution GAN outputs and established the architectural conventions still used today. The DCGAN paper's guidelines (no BN in  $D$  input layer, LReLU slope 0.2, Adam with  $\beta_1 = 0.5$ ) remain standard practice.

# Comparison of GAN Variants

| Variant     | Key change    | Loss        | Advantage     |
|-------------|---------------|-------------|---------------|
| Vanilla GAN | FC networks   | JS div.     | Simple        |
| DCGAN       | Conv. nets    | JS div.     | Better images |
| WGAN        | Wt. clipping  | Wasserstein | More stable   |
| WGAN-GP     | Grad. penalty | Wasser.+GP  | No clipping   |
| cGAN        | Conditioning  | JS or W     | Controlled    |
| StyleGAN    | Style-based   | W           | SOTA          |

# Key Equations: Summary

| Concept               | Formula                                                                                                |
|-----------------------|--------------------------------------------------------------------------------------------------------|
| Minimax objective     | $\min_G \max_D \mathbb{E}_{p_r}[\log D(\mathbf{x})] + \mathbb{E}_{p_z}[\log(1 - D(G(\mathbf{z})))]$    |
| Optimal discriminator | $D^*(\mathbf{x}) = p_r(\mathbf{x}) / (p_r(\mathbf{x}) + p_g(\mathbf{x}))$                              |
| Nash equilibrium      | $p_g = p_r, \quad D^* \equiv 1/2$                                                                      |
| Loss at equilibrium   | $V(G^*, D^*) = -2 \log 2$                                                                              |
| GAN loss = JS div.    | $V(G, D^*) = 2D_{\text{JS}}(p_r \  p_g) - 2 \log 2$                                                    |
| Non-sat. generator    | $\mathcal{L}_G = -\mathbb{E}_{p_z}[\log D(G(\mathbf{z}))]$                                             |
| Discriminator loss    | $\mathcal{L}_D = -\mathbb{E}_{p_r}[\log D(\mathbf{x})] - \mathbb{E}_{p_z}[\log(1 - D(G(\mathbf{z})))]$ |
| Wasserstein distance  | $W(p_r, p_g) = \sup_{\ f\ _L \leq 1} (\mathbb{E}_{p_r}[f] - \mathbb{E}_{p_g}[f])$                      |
| Gradient penalty      | $\lambda \mathbb{E}[(\ \nabla f_w(\hat{\mathbf{x}})\ _2 - 1)^2]$                                       |
| MMD                   | $\sqrt{\mathbb{E}_{p \times p}[k] - 2\mathbb{E}_{p \times q}[k] + \mathbb{E}_{q \times q}[k]}$         |
| FID                   | $\ \mu_r - \mu_g\ ^2 + \text{Tr}(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{1/2})$                    |

- 1 **Optimal discriminator.** Derive  $D^*(\mathbf{x}) = p_r/(p_r + p_g)$  by differentiating  $f(t) = A \log t + B \log(1 - t)$  and verify it is a maximum, not a minimum.
- 2 **JS divergence bounds.** Prove  $D_{\text{JS}}(p||q) \in [0, \log 2]$  for any two distributions and find the conditions for equality on each side.
- 3 **Vanishing gradient.** At  $D(G(\mathbf{z})) = \varepsilon \ll 1$ , compute  $|\nabla_{\theta(g)} \mathcal{L}_G^{\text{orig}}|$  and  $|\nabla_{\theta(g)} \mathcal{L}_G^{\text{non-sat}}|$  and compare their magnitudes as  $\varepsilon \rightarrow 0$ .
- 4 **Label smoothing.** Derive the optimal discriminator  $D_s^*(\mathbf{x})$  when real targets are smoothed to  $s \in (0, 1)$  instead of 1. Show it converges to  $D^*$  as  $s \rightarrow 1$ .
- 5 **Wasserstein distance.** Compute  $W(\mathcal{N}(0, 1), \mathcal{N}(\mu, 1))$  analytically and verify that  $W \rightarrow 0$  as  $\mu \rightarrow 0$ , unlike  $D_{\text{JS}}$  when the two Gaussians have non-overlapping support.
- 6 **Conditional GAN.** Extend the minimax objective to the conditional setting. What changes in the derivation of  $D^*(\mathbf{x}|\mathbf{y})$ ?

# Three Families: A Unified Mathematical View

All three learn  $p_{\theta}(x)$  but with different objectives and inductive biases.

| Criterion       | VAE            | Diffusion         | GAN          |
|-----------------|----------------|-------------------|--------------|
| Objective       | ELBO           | ELBO ( $T$ terms) | Minimax / JS |
| Likelihood      | ✓              | ✓                 | ×            |
| Latent space    | Compact $h$    | Full-dim $x_t$    | Noise $z$    |
| Encoder         | Learned        | Fixed schedule    | None         |
| Sampling cost   | 1 decoder pass | $T$ steps         | 1 pass       |
| Stability       | High           | High              | Low–Moderate |
| Sample quality  | Moderate       | State of the art  | High (sharp) |
| Mode coverage   | Good           | Excellent         | Partial      |
| Anomaly detect. | ✓              | ✓                 | ×            |



# Why VAEs Produce Blurry Samples

**Mathematical reason.** The VAE decoder minimises:

$$-\mathbb{E}_{q_{\phi}(\mathbf{h}|\mathbf{x})}[\log p_{\theta}(\mathbf{x}|\mathbf{h})] \propto \mathbb{E}_{q_{\phi}}[\|\mathbf{x} - \hat{\mathbf{x}}_{\theta}(\mathbf{h})\|^2].$$

Because the posterior  $q_{\phi}(\mathbf{h}|\mathbf{x})$  has nonzero variance, the decoder learns the **posterior mean over all likely completions**, which is a blur over the data manifold.

**Contrast with the other two models:**

## Diffusion avoids blurring by

predicting only the *noise added at one step*  $t$ , not the full image. Each reverse step makes a small, committed correction. The iterative nature removes the need to average over the full distribution of plausible outputs.

## GANs avoid blurring by

replacing the fixed  $\ell_2$  loss with an *adaptive, learned* discriminator loss. The discriminator penalises blurriness and any systematic deviation from real data, providing a richer training signal than any fixed pixel-wise objective.

# The Likelihood Problem in GANs

**A fundamental asymmetry.** VAEs and diffusion models optimise a lower bound on  $\log p(\mathbf{x})$ , so they can both *generate* and *evaluate* density. GANs have **no access to**  $p_g(\mathbf{x})$ .

## Consequences of no likelihood:

- Cannot compute log-likelihood scores for anomaly detection.
- Cannot evaluate model fit on held-out data in a principled way.
- Evaluation requires proxy metrics (FID, IS, MMD) rather than test-set log-likelihood.
- Cannot be embedded as a prior in a Bayesian framework without approximation.

**Why this is a price worth paying (sometimes).** The discriminator provides a *perceptual* loss that no fixed  $p_\theta$  objective can replicate: it adapts to what “looks real” to a neural network, capturing texture, sharpness, and high-frequency detail that  $\ell_2$  and even ELBO objectives systematically underweight.

# Training Stability: A Comparison

## VAE — Stable

ELBO is a well-defined scalar. Both networks trained with standard SGD. KL prevents collapse when weighted correctly. Main challenge: choosing  $\beta$  for  $\beta$ -VAE.

## Diffusion — Stable

MSE on noise prediction is smooth. No adversarial dynamics. Low MC gradient variance ( $x_t$  sampled in one step). Main challenge: noise schedule design.

## GAN — Unstable

Non-stationary minimax landscape. No Nash convergence guarantee.  $V(G, D)$  is concave-convex; simultaneous gradient ascent-descent diverges even for bilinear games. Requires WGAN-GP, spectral norm, label smoothing.

## Key insight

GAN instability is *fundamental*: the minimax saddle structure prevents convergence guarantees that both VAE and DDPM enjoy.

# Mode Collapse vs. Posterior Averaging: Two Failure Modes

The two pathological failure modes of generative models lie at opposite ends:

## Posterior averaging (VAE)

**Too much spread.** The decoder must explain all data through a single mean. Result: generated images are averages over the data manifold — blurry. The model *covers* all modes but fails to commit to any one of them.

$$\hat{x} = \mathbb{E}_{q_{\phi}(h|x)}[\hat{x}_{\theta}(h)] \approx \text{mean of modes}$$

## Mode collapse (GAN)

**Too little spread.** The generator finds a single safe output that fools  $D$  and ignores the rest of  $p_r$ . Result: sharp images but only from a few modes.  $\text{supp}(p_g) \ll \text{supp}(p_r)$ .

$$G(z) \approx x^* \forall z \Rightarrow D_{\text{JS}}(p_r \| p_g) \approx \log 2 \text{ (flat)}$$

**Diffusion avoids both:** the fixed encoder  $q(x_t|x_{t-1})$  prevents collapse (no network can “cheat” the noise schedule), and the iterative denoising avoids averaging because each step commits to only a small noise correction rather than a full image prediction.

## Pros

- **Likelihood bound:** ELBO usable for anomaly detection.
- **Compact latent:**  $\mathbf{h}$  supports interpolation, clustering, disentanglement, property prediction.
- **Fast generation:** single decoder pass.
- **Stable training:** well-defined ELBO, no adversary.
- **Flexible:**  $\beta$ -VAE for disentanglement; hierarchical VAEs for expressivity.
- **Scientific use:** ideal for molecules, spectra, simulations.

## Cons

- **Blurry samples:** posterior averaging is fundamental, not a capacity issue.
- **Posterior collapse:** decoder ignores  $\mathbf{h}$ ;  $\text{KL} \rightarrow 0$ , latent useless.
- **ELBO gap:**  
 $\log p(\mathbf{x}) = \text{ELBO} + D_{\text{KL}}(q_\phi \| p(\mathbf{h}|\mathbf{x})) > \text{ELBO}.$
- **Mean-field:** factorised Gaussian misses correlated posteriors.
- **Fixed-dim bottleneck:** wrong  $d_h$  hurts quality.

# Detailed Pros and Cons: Diffusion Models

## Pros

- **Best sample quality:** sharp, diverse, no blurring.
- **Full coverage:** fixed forward process prevents mode collapse.
- **Stable training:** MSE on noise, no adversarial dynamics.
- **Likelihood bound:**  $\text{ELBO} \rightarrow \log p(x)$  as  $T \rightarrow \infty$ .
- **Flexible conditioning:** text, class, depth via classifier-free guidance and cross-attention.
- **Rich theory:** SDE/score connection (Song et al. 2021).

## Cons

- **Slow sampling:**  $T=100$ – $1000$  steps (DDIM/DPM-Solver: 10–50).
- **No compact latent:** full-dim  $x_t$  precludes clustering or property prediction.
- **Compute-intensive:** large  $T$ , high-dim  $x$  need significant GPU memory.
- **Poor for repr. learning:** no encoder for features.
- **Schedule-sensitive:** cosine preferred over linear at high resolution.

# Detailed Pros and Cons: GANs

## Pros

- **Sharpest samples:** adversarial loss catches blurriness that  $\ell_2$ /ELBO miss.
- **Fastest inference:** single generator pass.
- **Flexible loss:** works when likelihood is intractable.
- **Rich ecosystem:** cGAN, DCGAN, WGAN-GP, StyleGAN, pix2pix.
- **Image translation:** handles paired and unpaired transfer.

## Cons

- **No likelihood:** no anomaly detection or density estimation.
- **Mode collapse:**  $G$  covers only part of  $p_r$ ; FID/IS may not detect it.
- **Instability:** non-stationary minimax; no convergence guarantee; hyperparameter-sensitive.
- **JS flatness:**  $D_{JS} = \log 2$  when supports don't overlap (WGAN fixes this).
- **Wasted discriminator:**  $D$  discarded after training.

# When to Use Each Model

| Use case                         | Model            | Reason                        |
|----------------------------------|------------------|-------------------------------|
| Image / audio / video synthesis  | Diffusion        | Best quality, no collapse     |
| Text-to-image (Stable Diffusion) | LDM              | VAE compression + diffusion   |
| Image-to-image translation       | GAN              | Adversarial perceptual loss   |
| Real-time generation             | VAE or GAN       | Single-pass inference         |
| Anomaly detection                | VAE or Diffusion | Likelihood bound available    |
| Representation / clustering      | VAE              | Compact structured latent     |
| Scientific data (molecules)      | VAE              | Low-dim latent for prediction |
| Data augmentation                | GAN or Diffusion | Sharp, diverse samples        |
| Inpainting / super-resolution    | Diffusion        | Score-based posterior         |
| Very limited data                | VAE              | Stable; KL regularisation     |



# Key Equations: Complete Summary Across All Three Models

| Concept              | VAE / Diffusion                                                                                                                                                                                                                  | GAN                                                                                     |
|----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|
| Encoder              | $q_{\phi}(\mathbf{h} \mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}_{\phi}, \boldsymbol{\sigma}_{\phi}^2 \mathbf{I})$<br>$q(\mathbf{x}_t \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{\alpha_t}\mathbf{x}_{t-1}, (1 - \alpha_t)\mathbf{I})$ | None (implicit via $G$ )                                                                |
| Marginal / forward   | $\mathbf{x}_t = \sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\boldsymbol{\epsilon}$                                                                                                                              | $\mathbf{x} = G(\mathbf{z}; \boldsymbol{\theta}^{(g)})$                                 |
| Objective            | $\mathcal{L} = \mathbb{E}[\log p_{\theta}(\mathbf{x} \mathbf{h})] - D_{\text{KL}}(q\ p)$                                                                                                                                         | $\min_G \max_D \mathbb{E}[\log D(\mathbf{x})] + \mathbb{E}[\log(1 - D(G(\mathbf{z})))]$ |
| Optimal solution     | DDPM: $\boldsymbol{\mu}_{\theta} = \frac{1}{\sqrt{\alpha_t}}(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}}\boldsymbol{\epsilon}_{\theta})$                                                                            | $D^*(\mathbf{x}) = \frac{p_r(\mathbf{x})}{p_r(\mathbf{x}) + p_g(\mathbf{x})}$           |
| Loss at equilibrium  | VAE ELBO $\rightarrow \log p(\mathbf{x})$                                                                                                                                                                                        | $V(G^*, D^*) = -2 \log 2$                                                               |
| Divergence minimised | $D_{\text{KL}}(q\ p)$ at each level                                                                                                                                                                                              | $D_{\text{JS}}(p_r\ p_g)$                                                               |
| Score connection     | $\mathbf{s}_{\theta} = -\boldsymbol{\epsilon}_{\theta} / \sqrt{1 - \bar{\alpha}_t}$                                                                                                                                              | $D^* \rightarrow$ implicit score estimator                                              |
| Sampling             | 1 pass (VAE); $T$ steps (DDPM)                                                                                                                                                                                                   | 1 pass                                                                                  |

# Exercises: Cross-Model Comparisons

## Exercises 1–3

- 1 **Common ELBO.** Show VAE (1 level) and DDPM ( $T$  levels) are special cases of the MHVAE ELBO. Identify encoder, decoder, prior in each.
- 2 **Blurriness from  $\ell_2$ .** For  $p_r = \frac{1}{2}\delta(x-a) + \frac{1}{2}\delta(x-b)$ , show the VAE decoder learns  $\hat{x} = (a+b)/2$  while a GAN can produce both modes.
- 3 **Convergence.** State which property of the DDPM ELBO makes it amenable to SGD while the GAN minimax saddle is not.

## Exercises 4–6

- 4 **Wasserstein vs. JS.** Compute  $D_{\text{JS}}(\mathcal{N}(0, 1) \parallel \mathcal{N}(\mu, 1))$  and  $W(\mathcal{N}(0, 1), \mathcal{N}(\mu, 1))$ . Show  $D_{\text{JS}}$  saturates at  $\log 2$  while  $W = |\mu|$  grows.
- 5 **Score matching bridge.** Using  $\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t | \mathbf{x}_0) = -\epsilon / \sqrt{1 - \bar{\alpha}_t}$ , show DDPM equals denoising score matching (Vincent 2011).
- 6 **Latent Diffusion.** Express the LDM objective as a composition of VAE ELBO and DDPM ELBO. Identify which parameters are trained in each stage.

## Diffusion models

- ❶ Kingma & Welling, *Auto-Encoding Variational Bayes*, ICLR 2014. 1312.6114
- ❷ Ho, Jain & Abbeel, *DDPM*, NeurIPS 2020. 2006.11239
- ❸ Song et al., *Score-Based Generative Modeling through SDEs*, ICLR 2021. 2011.13456
- ❹ Song, Meng & Ermon, *DDIM*, ICLR 2021. 2010.02502
- ❺ Kingma et al., *Variational Diffusion Models*, NeurIPS 2021. 2107.00630
- ❻ Nichol & Dhariwal, *Improved DDPM*, ICML 2021. 2102.09672
- ❼ Rombach et al., *Latent Diffusion Models*, CVPR 2022. 2112.10752
- ❽ Luo, *Understanding Diffusion Models*, 2022. 2208.11970

## Generative Adversarial Networks

- ❶ Goodfellow et al., *Generative Adversarial Nets*, NeurIPS 2014. 1406.2661
- ❷ Mirza & Osindero, *Conditional GAN*, 2014. 1411.1784
- ❸ Radford, Metz & Chintala, *DCGAN*, ICLR 2016. 1511.06434
- ❹ Arjovsky, Chintala & Bottou, *WGAN*, ICML 2017. 1701.07875
- ❺ Gulrajani et al., *WGAN-GP*, NeurIPS 2017. 1704.00028
- ❻ Heusel et al., *FID*, NeurIPS 2017. 1706.08500
- ❼ Goodfellow, Bengio & Courville, *Deep Learning*, MIT Press 2016, Ch. 20.
- ❽ Weng, *From GAN to WGAN*, blog post, 2017.